
多线程
Multithreading

现代C++基础 Modern C++ Basics

Jiaming Liang, undergraduate from Peking University

- **Thread**
- **Synchronization utilities**
- **High-level Abstraction of Asynchronous Operation**

Multithreading

Thread

Multithreading

- Thread
 - Abstract thread model
 - thread
 - jthread
 - Miscellaneous topics

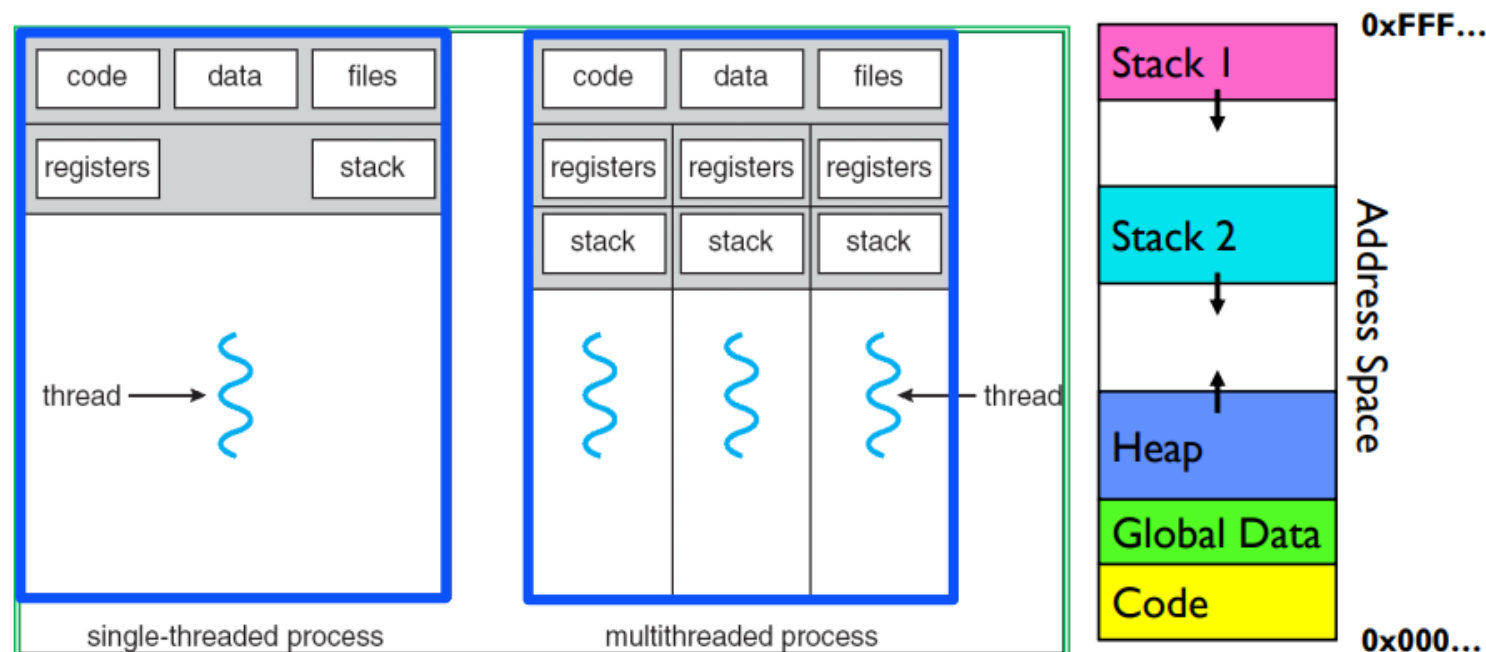
Thread model

- We first briefly review what thread is...
- We've learnt in ICS that each program is a **process**;
 - It has independent address space, and possibly other status like file descriptors (depend on OS).
 - Good isolation, good protection, really limited ability to access memory of another process.
- Threads: less protection, better data sharing!
 - They still partially keep their own set of resources (like registers)...
 - But lie in the same virtual address space, so easy to access memory of other threads!
 - Usually, OS will schedule threads instead of processes;
 - So to some extent, we can say threads are the smallest units to utilize multi-core parallelism.

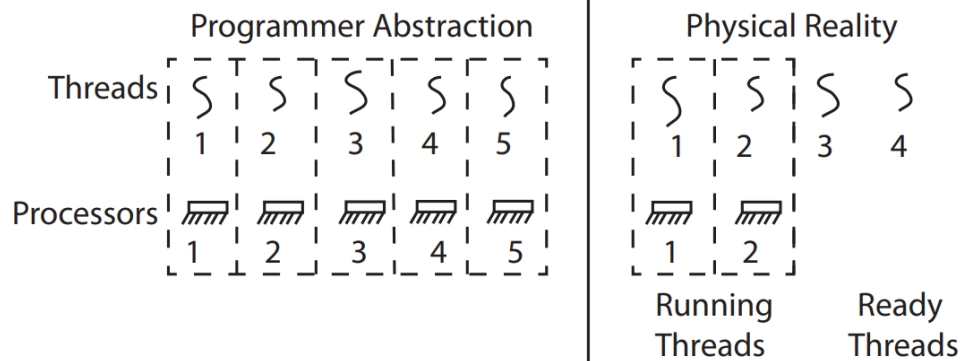
BTW: C++ standard in fact doesn't have a concept of "process", since all things it regulates happen in a single program; you need OS-dependent APIs to manipulate processes, like [fork/exec](#) in Linux, or external libraries like Boost.Process.

Thread model

- In a nutshell:



Roughly speaking, if there are two physical cores, there are actually only two threads executing simultaneously; but OS gives an illusion that more than two threads run concurrently by scheduling.



Credit: Prof. Jin Xin @ PKU OS.

We only give a very rough understanding of thread; you need to learn it comprehensively in OS course.

Thread model

- Scheduling is in fact pausing a running thread, then running a ready thread.
 - And scheduling algorithms determine which to pause and which to run.
 - Registers will be saved and restored during context switch.
- Threads compete with each other for executing themselves!
 - Thus, you may think statements may execute in any order, which leads to data race and synchronization problems.^[1]
- Threads need to be **joined** or **detached** after creation; the former will **wait** until the thread function exits, and the latter will make it execute separately and freely.

[1]: We'll give a rigorous definition for data races in *Advanced Concurrency*.

Multithreading

- Thread
 - Thread model in computers
 - thread
 - jthread
 - Miscellaneous topics

Thread

- We've learnt in ICS how to use **pthread** in POSIX system.

- **pthread_create/join/...**, like this:

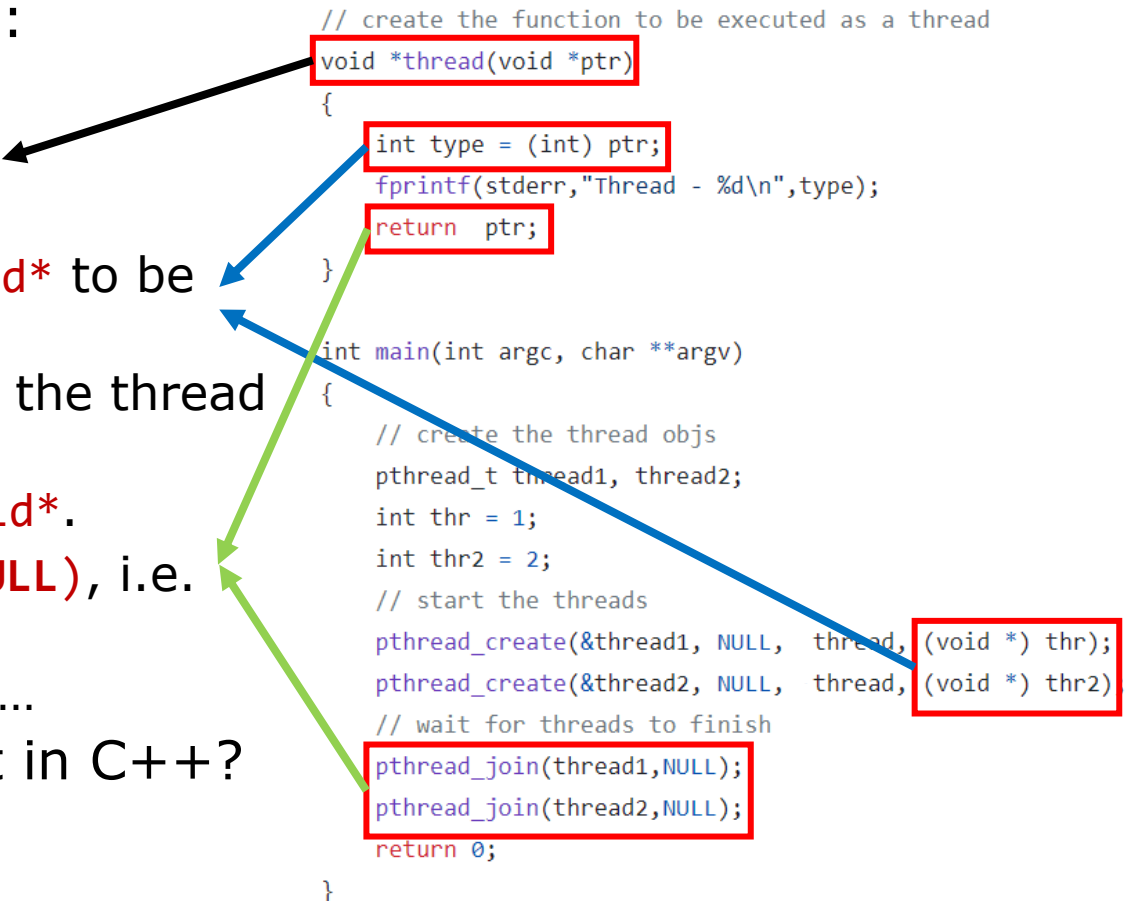
- Obscure C interface...

1. The thread function should be **void* func(void*)**;
2. Parameters are packed in a **void*** to be passed in.
 - You need to unpack it inside the thread function, like **(int)** here.
3. Return value is accepted by **void***.
 - Here we **pthread_join(..., NULL)**, i.e. not need return value.
 - If you need it, unpack again...

- Very strange... can we improve it in C++?

```
// create the function to be executed as a thread
void *thread(void *ptr)
{
    int type = (int) ptr;
    fprintf(stderr, "Thread - %d\n", type);
    return ptr;
}

int main(int argc, char **argv)
{
    // create the thread objs
    pthread_t thread1, thread2;
    int thr = 1;
    int thr2 = 2;
    // start the threads
    pthread_create(&thread1, NULL, thread, (void *) thr);
    pthread_create(&thread2, NULL, thread, (void *) thr2);
    // wait for threads to finish
    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);
    return 0;
}
```



Thread

- Of course! We may code like:
- Very intuitive, very simple, by `std::thread` defined in `<thread>`.
 1. It isn't limited to a function, but can use any functor.
 - It's equivalent to call `std::invoke`, so you can also pass into pointer to member function with `this`, like `{ &SomeClass::MemberFunc, this, params... }`;
 - Functor will be copied; you can explicitly `std::move(funcutor)` to move it or `std::ref` to make it a `std::reference_wrapper`.
 2. You can call `.detach()` to detach the thread.
 - After join/detach, the underlying thread is not associated with `std::thread` object. Unless you move a new object to it, this `std::thread` object is in an empty state (just like default-constructed/moved).
 3. Return value of `func` will be omitted; it should be passed by ref. param.

```
void func(int num)
{
    std::println("Thread info: {}", num);
}

int main()
{
    std::thread t{ func, 1 };
    t.join();
}
```

Thread

- And some other APIs:
 1. Move ctor, move assignment, and swappable (by `.swap()` or `std::swap`).
 2. `.joinable() -> bool`: whether the thread is in an empty state.
 - E.g. before calling `.join()/.detach()`, it returns `true`; afterwards `false`.
 3. Dtor: `std::terminate` if `.joinable()`, otherwise do nothing;
 - That is, every running `std::thread` must call `.join()/.detach()` before destruction.
 4. `.get_id()`: get thread id;
 - Class `std::thread::id` instead of simply an integer.
 - Restricted integer: only comparable, hashable and printable (by `<<` since C++11 or `std::formatter` since C++23).
 - Particularly, it only supports `fill-align-width` format.
 - `std::thread` in empty state will get default-constructed `id`.

Thread

```
INVOKE( decay-copy( std::forward<F>(f) ),  
        decay-copy( std::forward<Args>(args) ) ... ) (until C++23)
```

```
std::invoke( auto( std::forward<F>(f) ),  
            auto( std::forward<Args>(args) ) ... ) (since C++23)
```

- **Note:** parameters are decay-copied to the thread functions.
 - And since C++23 it can be explained by **auto(...)** + materialization, i.e. generate a prvalue that is materialized in the current thread.
 - So **auto** will decay type, e.g. **Object& -> Object**.
 - When forwarded type is **Object&**, then copy;
 - When **Object&&**, then move;
- Let's break it down step by step...

```
void func(Object object)
{
    std::cout << "Thread id: " << std::this_thread::get_id() << "\n";
}

int main()
{
    std::cout << "Main id: " << std::this_thread::get_id() << "\n";
    std::thread t{ func, Object{} };
    t.join();
}
```

```
class Object
{
public:
    Object() { std::cout << "Construct at " << std::this_thread::get_id() << "\n"; };
    ~Object() { std::cout << "Destruct at " << std::this_thread::get_id() << "\n"; };
    Object(const Object&) {
        std::cout << "Const Copy at " << std::this_thread::get_id() << "\n";
    };
    Object(Object&&) { std::cout << "Move at " << std::this_thread::get_id() << "\n"; };
    Object& operator=(const Object&) {
        std::cout << "Const Copy Assignment at " << std::this_thread::get_id() << "\n";
        return *this;
    };
    Object& operator=(Object&&) {
        std::cout << "Move Assignment at " << std::this_thread::get_id() << "\n";
        return *this;
    };
};
```

Thread

```
std::thread t{ func, Object{} };
```

```
template< class F, class... Args >  
explicit thread( F&& f, Args&&... args );
```

1. Parameters are passed into ctor of `std::thread`.
 - Encountering reference, prvalue `Object{}` is materialized and thus constructed as `arg0`.

2. `arg0` is then decay-copied.

- Here it's prvalue `Object{ std::move(arg0) }`, and materialized (and thus move-constructed).

```
INVOKE(decay-copy(std::forward<F>(f)),  
        decay-copy(std::forward<Args>(args))...) (until C++23)
```

```
std::invoke(auto(std::forward<F>(f)),  
            auto(std::forward<Args>(args))...) (since C++23)
```

3. Then thread executes `std::invoke`;

- Materialized parameters are passed to new thread.

```
template< class F, class... Args >  
std::invoke_result_t<F, Args...>  
invoke( F&& f, Args&&... args )
```

4. And finally `arg1` is forwarded to parameters of `func`.

- And thus move-constructed as `object`.

1) Invoke the *Callable* object `f` with the parameters `args` as by

```
INVOKE(std::forward<F>(f), std::forward<Args>(args)...) (since C++23)
```

```
void func(Object object)  
{  
    ...  
    std::cout << "Thread id: " << std::this_thread::get_id() << "\n";  
}
```

Thread

You'll implement `std::thread` yourself in our homework to know how these really happen.

- The first three steps happen at the current thread, and the final step happens at the new thread.
- So the output is like:
 - Any exception thrown in step 1 & 2 can then be caught in old thread.
- Exercise: is this piece of code right?
 - No, since we forward materialized temporary to `func`, i.e. `func(std::move(...))`;
 - And you cannot bind rvalue to lvalue reference...
 - Even if you use `const int&`, it in fact refers to a temporary, not the parameter you passed to thread ctor.

```
Main id: 21244
Construct at 21244
Move at 21244
Destruct at 21244
Move at 8504
Thread id: 8504
Destruct at 8504
Destruct at 8504
```

Step 1
Step 2; new thread executes Step 3.
Param of thread ctor destructed.
Step 4

Param of func destructed.
Materialized temporary destructed.

```
#include <thread>
#include <iostream>

void func(int& type) {
    type = 2;
}

int main()
{
    int type = 1;
    std::thread t{ func, type };
    t.join();
    return 0;
}
```

```
/opt/compiler-explorer/gcc-13.2.0/include/c++/13.2.0/bits/std_thread.h:157:72: error: static
assertion failed: std::thread arguments must be invocable after conversion to rvalues
157 |                                     typename decay<_Args>::type...>::value,
    |                                     ^~~~~~
```


Thread

- Reason: decay-copy instead of reference is safer.
 - We've learnt in ICS that you may pass a pointer to another thread, so that another thread can access memory of current thread.
 - If the referred object goes out of its lifetime, you're accessing invalid memory!
 - Simultaneous access in different threads may lead to data races too.
- You need to use `std::(c)ref()` explicitly to pass the (const) reference.
 - You've seen similar way in `std::bind_xx`, which also warns you about lifetime problem.

```
std::thread t{ func, std::ref(type) };  
t.join();  
std::cout << type; 2请按任意键继续. . .
```

Thread

```
native_handle_type native_handle();
```

(since C++11)
(not always present)

Returns the implementation defined underlying thread handle.

- Note 1: APIs provided in `std::thread` are high-level; sometimes you may want fine-grained control.
 - For example, you may want to change the priority of some threads.
 - It's platform-dependent, so C++ provides a `.native_handle()` for it.
 - The return type is platform-dependent (e.g. `pthread_t` in POSIX system).
- Note 2: `static constexpr std::thread::hardware_concurrency()` can be used to check number of real parallel threads.
 - Roughly speaking, how many physical cores.
 - This is only a hint to the possible maximum parallelism; you need profiling to get the best thread number for your program's performance.
 - When the system cannot give a hint, return 0.

Thread

- Note 3: some rare but possible exceptions, listed here.

- All exceptions are `std::system_error` with some error code.
- Ctor: *Throws:* `system_error` if unable to start the new thread.

Error conditions: i.e. error code

— `resource_unavailable_try_again` — the system lacked the necessary resources to create another thread, or the system-imposed limit on the number of threads in a process would be exceeded.

- `.join()/.detach()`:

Error conditions:

— `resource_deadlock_would_occur` — if deadlock is detected or `get_id() == this_thread::get_id()`. A thread waiting for itself; `.detach()` doesn't have this case.

— `no_such_process` — if the thread is not valid.

— `invalid_argument` — if the thread is not joinable.

Thread

- Note 4: `namespace std::this_thread` has many methods for the current thread.
 - `get_id()`: get id of current thread.
 - `sleep_for()/sleep_until()`: pause the current thread.
 - `yield()`: request scheduling.
 - We've said that threads compete with each other; OS will schedule a thread when it has executing for a period of time.
 - That is, a thread will execute eagerly, and OS forces it to pause.
 - `yield` means the thread gives up execution right voluntarily, and OS re-schedules it.
 - However, OS may still choose the original thread to run, if the priority of this thread is high enough so that scheduling algorithms still choose it.
 - i.e. pause the thread, save its state, and reload the same state, and continue to run.

Multithreading

- Thread
 - Thread model in computers
 - thread
 - jthread
 - Miscellaneous topics

jthread

- C++ encourages RAII, which means dtor will release resource acquired by ctor.
 - But `std::thread` seems to violate it, because if you forget to join/detach a thread, then the whole program is terminated.
 - `std::jthread` in `<thread>` since C++20 is used to solve that; it will automatically join the thread if its `joinable()` is still `true` in dtor.
 - It has all APIs of `std::thread`, i.e. you can use `join/detach/swap/move/native_handle/joinable/get_id/hardware_concurrency`.
 - But move will try to join thread it holds, and then move another to `*this`.
 - **Self-move will also join itself!**
- In C++11, some argue that termination is better than silent wait, which makes it not default behavior for `std::thread`.

From the start (pre-C++11), many (incl. me) had wanted `thread` to have what is now `jthread`'s behavior, but people grounded in traditional operating systems threads insisted that terminating a program was far preferable to a deadlock. In 2012 and 2013, Herb Sutter proposed a joining

jthread

- Besides, `std::jthread` also adds **stop token handling**.
 - Also sometimes called *cooperative cancellation*.
 - We know that threads compete and execute eagerly.
 - You can seldom *force* a thread to do something, but **request** it to do.
 - For stopping a thread, you may set some shared data, and the thread checks it periodically; when check succeeds, it **returns voluntarily**.
-
- Note that you can NOT kill a thread (though you can kill a process), since data dependence in threads is too common.
 - For example, what if a thread is still holding a lock, but it's forced to exit? Then the waiting thread will go into deadlock!
 - All in all, you can hardly ever guarantee a thread to be in a consistent state when you kill it; that's why we need stop token as a hint.

Stop token

- So the requester holds a ***stop source***, and the thread holds a ***stop token*** that associates with the stop source.
 - To prevent use-after-free, they share an underlying *stop state* with reference counts;
 - The state records related information and will be freed when counts goes to 0.
 - The stop source can only request once, which sets some flag in state;
 - Future requests have no actual effects.
 - And the stop token can check regularly whether the flag is set.
- So the state should expose interface below:
 - Setter: `request_stop()`, set the flag;
 - Getter: `stop_requested()`, check whether the flag is set;
 - Share/Detach: increment/decrement reference count.

Stop token

- And accordingly, `std::stop_source` and `std::stop_token` in `<stop_token>` wrap and expose them in a thread-safe way.
- `std::stop_source`:
 - Setter: `.request_stop()`;
 - Getter: `.stop_requested()`;
 - Share & Detach:
 - Default ctor: create a `stop_source` with newly created state.
 - Copy ctor & assignment: share the current state with others;
 - Move ctor & assignment: transfer the ownership;
 - Dtor: detach.
- `.get_token()` -> `std::stop_token`: get a stop token that shares the same state.

Stop token

- `std::stop_token`:
 - No setter;
 - Getter: `.stop_requested()`;
 - Share & Detach:
 - Copy ctor & assignment: share the current state with others;
 - Move ctor & assignment: transfer the ownership;
 - Dtor: detach.
- For example:

```
void func(std::stop_token token, int& cnt)
{
    while (!token.stop_requested())
    {
        cnt++;
    }
}
```

```
int main()
{
    int cnt = 0;
    using namespace std::literals; // To use literal suffix 'ms'.

    std::stop_source source;
    std::thread t{ func, source.get_token(), std::ref(cnt) };
    std::this_thread::sleep_for(1ms);
    source.request_stop();
    t.join();

    return 0;
}
```


Stop token

- They can also attach to an empty state, then every operation does nothing.
 - `std::stop_token`: default construct it;
 - `std::stop_source`: add a placeholder tag:

```
explicit stop_source( std::nostopstate_t nss ) noexcept;    (2)  std::stop_source source{ std::nostopstate };
```

- Since default ctor will create a new state.
 - When they're e.g. moved-from, then the state will be empty too.
 - To check whether the current state is empty, you can use method `.stop_possible()`; it returns `false` when empty.
 - And `.stop_requested()` also returns `false` when empty.
- Particularly, when only stop tokens associate with a state that hasn't been requested (i.e. no stop source, so no possible request), `token.stop_possible()` also returns `false`.

- To conclude:

std::stop_token

Member functions

(constructor)	constructs new stop_token object (public member function)
(destructor)	destructs the stop_token object (public member function)
operator=	assigns the stop_token object (public member function)

Modifiers

swap	swaps two stop_token objects (public member function)
------	--

Observers

stop_requested	checks whether the associated stop-state has been requested to stop (public member function)
stop_possible	checks whether associated stop-state can be requested to stop (public member function)

Non-member functions

operator== (C++20)	compares two std::stop_token objects (function)
swap (std::stop_token) (C++20)	specializes the std::swap algorithm (function)

std::stop_source

Member functions

(constructor)	constructs new stop_source object (public member function)
(destructor)	destructs the stop_source object (public member function)
operator=	assigns the stop_source object (public member function)

Modifiers

request_stop	makes a stop request for the associated stop-state, if any (public member function)
swap	swaps two stop_source objects (public member function)

Observers

get_token	returns a stop_token for the associated stop-state (public member function)
stop_requested	checks whether the associated stop-state has been requested to stop (public member function)
stop_possible	checks whether associated stop-state can be requested to stop (public member function)

Non-member functions

operator== (C++20)	compares two std::stop_source objects (function)
swap (std::stop_source) (C++20)	specializes the std::swap algorithm (function)

Stop token

- So how does `std::jthread` cooperate with stop token?
 1. It contains a default-constructed `std::stop_source` directly.
 - `.get_stop_source()` to get a copy;
 - `.get_stop_token()` to get a token associated with underlying source;
 - Equivalent to `underlying_source.get_token()`.
 - And `.request_stop()`, equivalent to `underlying_source.request_stop()`.
 2. In dtor of `std::jthread`, if the source hasn't issued a request, call `.request_stop()`;
 - RAII to some extent.
 3. When possible, it will pass `get_token()` to its functor.

The new thread of execution starts executing:

```
std::invoke(auto(std::forward<F>(f)), get_stop_token(),  
            auto(std::forward<Args>(args))...) (since C++23)
```

if the expression above is well-formed, otherwise starts executing:

```
std::invoke(auto(std::forward<F>(f)),  
            auto(std::forward<Args>(args))...) (since C++23)
```

Stop token

- For example:

`get_stop_token()` is provided automatically.

```
using namespace std::literals;
std::jthread t{ [](std::stop_token token) {
    while (!token.stop_requested())
    {
        std::cout << "PKU No.1!";
    }
} };
std::this_thread::sleep_for(1s);
t.request_stop();
```

This may be omitted since
dtor of `std::jthread` will
automatically `.request_stop()`.

Note that `std::cout` is safe to be used in multiple threads; other streams should use `std::osyncstream` in C++20.

Stop token

- Another example:

```
using namespace std::literals;
std::jthread t{ [](std::stop_token token) {
    while (!token.stop_requested())
    {
        std::cout << "PKU No.1!\n";
    }
} };

std::jthread t2{ [](std::stop_token token) {
    while (!token.stop_requested())
    {
        std::cout << "THU No.2!\n";
    }
}, t.get_stop_token() };

std::this_thread::sleep_for(1s);
t.request_stop();
```

Question: can we omit `t.request_stop()` here?

No! Since `t2` is destructed first, so `t2.join()` is before `t.request_stop()` in dtor of `t`.

Thus infinite loop in `t2`...

`t2` doesn't use its own `.get_token()` in functor; the functor shares the same state with `t`.

The new thread of execution starts executing:

```
std::invoke(auto(std::forward<F>(f)), get_stop_token(),
            auto(std::forward<Args>(args))...) (since C++23)
```

if the expression above is well-formed, otherwise starts executing:

```
std::invoke(auto(std::forward<F>(f)),
            auto(std::forward<Args>(args))...). (since C++23)
```

Stop token

- Finally, stop request can be associated with callbacks.

- By `std::stop_callback` with `std::stop_token`:

- Ctor registers callback on the state;

```
template< class C >
explicit stop_callback( const std::stop_token& st, C&& cb ) noexcept(/*see below*/);
```

- Dtor deregisters the callback.

- For example:

Callbacks will be executed here.

```
using namespace std::literals;
std::jthread t{ [](std::stop_token token) {
    while (!token.stop_requested())
    {
        std::cout << "PKU No.1!\n";
    }
} };

std::stop_callback callback{
    t.get_stop_token(),
    []() { std::cout << "THU No.2!\n"; }
};

std::this_thread::sleep_for(1s);
t.request_stop();
```

Stop token

- Note 1: it's quite like doing callback in a thread-safe way.
 - Callbacks will be executed exactly once for multiple requests;
 - Register and deregister are thread-safe; Deregister in a thread will wait for invocation in another thread if they happen in parallel.
- Note 2: the thread that first calls `.request_stop()` will execute all callbacks;
 - If there are multiple callbacks, the execution order is not regulated.
- Note 3: when request has been issued before registering...
 - i.e. in ctor of `std::stop_callback`, `token.stop_requested() == true`;
 - Then callback will be executed immediately in ctor in the current thread.
- Note 4: callback is not allowed to throw; `std::terminate()` if exception is thrown out of callback (treated as if `noexcept`).

Multithreading

- Thread
 - Thread model in computers
 - thread
 - jthread
 - Miscellaneous topics

Exception in threads

- We know that if we don't catch an exception in a single-threaded program, then `std::terminate`.
- Generally speaking, any thread that doesn't catch its exception when exiting will lead to `std::terminate`.
- So code right doesn't work:

```
void func()
{
    throw std::runtime_error{ "Not implemented" };
}

int main()
{
    try {
        std::thread t{ func }; // std::terminate!
        t.join();
    }
    catch (const std::runtime_error& err)
    {
        // ...
    }
}
```

Exception in threads

- So how can we pass the exception out of the thread?
- By `std::exception_ptr` defined in `<exception>`!
 - Roughly speaking, it's a shared pointer to exception.
 - Only when all pointers to the exception object destruct will the object destruct.
 - The actual type is implementation-defined... `using exception_ptr = /*unspecified*/` (since C++11)
 - It's regulated to expose these interfaces:
 - Default construct: as if it's a `nullptr`;
 - `std::make_exception_ptr(Exception)`: make a pointer that copies `Exception`;
 - And can be converted to `bool`, like a pointer.
 - `std::current_exception()`: used in catch block, as if `make_exception_ptr` to the current caught exception;
 - `std::rethrow_exception(std::exception_ptr)`: rethrow the exception object.

Exception in threads

- For example:

Pass exception
out of thread
by parameter.

```
void THUStudent() {  
    throw std::runtime_error("THU is not best!");  
    std::cout << "THU is best.\n";  
}  
  
void PKUStudent() {  
    std::cout << "PKU is best.\n";  
}  
  
void Work(std::exception_ptr& ptr)  
{  
    try {  
        PKUStudent();  
        THUStudent();  
    }  
    catch (const std::runtime_error&) {  
        ptr = std::current_exception();  
    }  
    return;  
}  
  
void Watch()  
{  
    std::exception_ptr ptr;  
    // join immediately.  
    {std::jthread _{ Work, std::ref(ptr) }; }  
    if (ptr)  
        std::rethrow_exception(ptr);  
    std::cout << "All students over.\n";  
}  
  
int main()  
{  
    try {  
        Watch();  
    }  
    catch (const std::runtime_error& error) {  
        std::cout << error.what();  
    }  
    return 0;  
}
```

Continue to throw
the exception in
the main thread.

PKU is best.
THU is not best!

Static block variable

- Previously we may use static variables in function to share it across calls.
 - But is it safe to use in multiple threads?
- Yes and no...
 - Yes: only one thread will execute initialization and other threads will wait (since C++11).
 - No: to modify it across multiple threads, you still need lock.

```
void Foo(int id) {  
    // Thread-safe: initialization will be executed exactly once.  
    static std::map<int, int> lookupTable{};  
    // Not thread-safe, need lock protection.  
    lookupTable.emplace(1, 2);  
}
```

Once-for-all

- More generally, if we want to execute some segment of code only once across all threads...

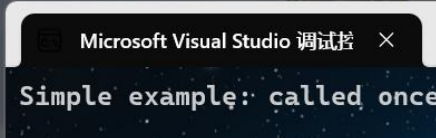
1. Tricks by static variable:

- This trick is often used in single-thread program too...

```
[[maybe_unused]] static int _ = []() {  
    HostUtils::CheckOptixError(optixInit());  
    return 0;  
}();
```

2. By `std::call_once` and `std::once_flag` defined in `<mutex>`:

```
std::once_flag flag1;  
  
void simple_do_once()  
{  
    std::call_once(flag1, []() { std::cout << "Simple example: called once\n"; });  
}  
  
int main()  
{  
    std::jthread st1{ simple_do_once }, st2{ simple_do_once },  
        st3{ simple_do_once }, st4{ simple_do_once };  
}
```



Once-for-all

Defined in header `<mutex>`

```
template< class Callable, class... Args >  
void call_once( std::once_flag& flag, Callable&& f, Args&&... args );
```

- `std::once_flag` only has a default ctor, meaning “not already called”.
 - And `std::call_once` will set the flag; exactly one thread will execute the callable and others will wait until it has completed.
- What’s the difference?
 1. `std::call_once` is slightly more flexible and intuitive;
 - But return value is ignored.
 2. Static variable trick may have slightly better performance.
 - See [stackoverflow](https://stackoverflow.com/questions/1649190/recursive-initialization-of-static-variables-is-undefined-behavior) for details.
 3. Recursive initialization for static variable is UB;
 - While `std::call_once` will lead to deadlock.

```
int Foo(int a, int b)  
{  
    static int m = Foo(a + 1, b + 1);  
    return m + 1;  
}
```

Once-for-all

- Sometimes we only want to share variables in calls of the current thread;
 - E.g. each thread has its own “static block variable”.
- We can use `thread_local` to specify thread storage duration!

This `static` can be omitted; `thread_local` block variables imply `static` if not specified. [See C++ Standard.](#)

```
void Foo(int id) {  
    ....  
    // Each thread has its own lookupTable.  
    thread_local static std::map<int, int> lookupTable{};  
    lookupTable.emplace(1, 2);  
    lookupTable.find(id);  
}
```

- Of course, you can use in “global” variables...

```
thread_local int m = 0; // external linkage  
static thread_local int n = 0; // internal linkage  
class A  
{  
    static thread_local int k;  
};  
thread_local int A::k = 0;
```

Note that it's not allowed to write `static` here.

Once-for-all

- `thread_local` global variables are created after a thread starts, and destructed when it exits.
- Final word: if an exception throws out of:
 - Initialization of static block variables;
 - Function of `std::call_once`;
- Then it's seen as execution failure, and it will be initialized / executed again the next time.
 - So pay attention if there are other side effects that cannot be executed twice.