

补充与总结

Supplementary and Summary

现代C++基础

Modern C++ Basics

Jiaming Liang, undergraduate from Peking University

Postgraduate from PKU since 2024.9 :-)

- **File system**
- **Time utilities (Chrono)**
- **Math utilities**
- **Summary and future prospect**

Supplementary and Summary File system

Overview

- Files represent organized data in non-volatile storage to let programs share data across different runs.
 - Files are named collection of data.
- Directories are used to construct file hierarchy.
 - Directories are named collection of files and directories.
- File system is an abstraction layer provided by OS to enable users to use *path* to access files and directories.
 - It records metadata of files and directories (size, modification time, owner, etc.) to make them organized and hierarchical.
- C++17 includes related utilities in `<filesystem>`.

Supplementary

- File system
 - Path operations
 - Overview
 - `std::filesystem::path`
 - File system operations

Path Overview

- Essentially, path is a string that represents location of a file.
- There are two different kinds of paths:
 1. Absolute path: always refer to the same location.
 2. Relative path: the location relative to *current working directory* (CWD) for the current process.
 - By changing CWD, the process can get different locations.
- A path consists of these components:
 1. Root name (optional): like drive name in Windows (C:, D:); or [UNC](#) ([//machine](#)), [etc.](#)
 2. Root directory (optional): a directory separator (\ on Windows, / on Linux, : on [classic MacOS](#)).
 3. Relative path: a sequence of filenames separated by directory separator.

Path Overview

- Besides no separators, filenames have many other platform-dependent characteristics or restrictions.

- (1.1) — The permitted characters. e.g. on Windows:

文件名不能包含下列任何字符:
\\/*?"<>|

[Example 1: Some operating systems prohibit the ASCII control characters (0x00 – 0x1F) in filenames. — end example]

[Note 1: Wider portability can be achieved by limiting *filename* characters to the POSIX Portable Filename Character Set:

A B C D E F G H I J K L M N O P Q R S T U V W X Y Z
a b c d e f g h i j k l m n o p q r s t u v w x y z
0 1 2 3 4 5 6 7 8 9 . _ — end note]

- (1.2) — The maximum permitted length. e.g. well-known [260](#) in some Windows functions.
 - (1.3) — Filenames that are not permitted. e.g. [CON on Windows](#).
 - (1.4) — Filenames that have special meaning.
 - (1.5) — Case awareness and sensitivity during path resolution. e.g. Linux is case-sensitive while Windows not.
 - (1.6) — Special rules that may apply to file types other than regular files, such as directories.

Particularly, . and .. means current directory and parent directory respectively.

Path Overview

- To make program cross-platform, C++ regulates a “generic format” that uses POSIX convention.
 - That is, these three components are just concatenated to form a path.
 - And `/` is considered as universal separator.
- Besides, C++ allows “native format” that depends on file system.
 - E.g. [OpenVMS](#), a legacy system previously used in bank (well, if it really supports C++17 utilities).

What Is a Fully Qualified Name?

A *fully qualified name* indicates how a file fits into a structure (a system of directories and subdirectories) that contains all the files stored under the OpenVMS system. The following type of file specification is a fully qualified name:

node::device:[directory]filename.file-type;version e.g. DKA0:[JDOE.DATA]test.txt

Path Overview

- Note 1: “generic” only means it’s a valid format in all systems; but its location is not always the same.
 - What is `D:\sub\path`?
 - Windows: an absolute path at D drive, with two components `sub` and `path`.
 - Linux: a relative path with name “`D:\sub\path`”, i.e. the whole string is a single component.
 - What is `/home/user`?
 - Linux: an absolute path with two components `home` and `user`.
 - Windows: a relative path at drive of CWD, with two components `home` and `user`.
- Note 2: it’s not very safe to rely on relative path since it’s just like a global variable, which can be changed by other threads and external library.
 - i.e. the location of a relative path can be modified arbitrarily.

Path Overview

- Note 3: “relative” or “absolute” only means whether it’s interfered by CWD; there can be many paths that refer to the same location.
 - E.g. `/home/user`, `/home/user/.`, `/home/user/dir/..`,
 - To unify all representations, we can do *path normalization*.
 - There are two kinds of normalization:
 - Lexical: a string-level substitution, which doesn’t change whether the path is relative or absolute.
 - So `/home/user`, `/home/user/.`, `/home/user/dir/..` are all normalized to `/home/user`, and paths `./user`, `./user/..`, `./user/dir/..` are all normalized to `user`.
 - Filesystem-dependent: paths are normalized to a unique absolute path.
 - Assuming CWD is `/home`, `./user`, `./user/..`, `./user/dir/..` are all normalized to `/home/user`.
 - C++ call such normalization as “canonical”.

Path Overview

- Specifically, a normalized path requires:

A path might be or can become normalized. In a normalized path:

- File names are separated only by a single preferred directory separator.
- The file name ". ." is not used unless the whole path is nothing but ". ." (representing the current directory).
- The file name does not contain "... ." file names (we do not go down and then up again) unless they are at the beginning of a relative path.
- The path only ends with a directory separator if the trailing file name is a directory with a name other than ". ." or "... .".

Path	POSIX normalized	Windows normalized
foo/././bar/.../	foo/	foo\
//host/.../foo.txt	//host/foo.txt	\\\host\foo.txt (Due to UNC path)
./f/.../.f/	.f/	.f\
C:bar/.../	.	C:
C:/bar/..	C:/	C:\
C:\bar\.../	C:\bar\...	C:\
./.../data.txt	/data.txt	\data.txt
././	.	.

Supplementary

- File system
 - Path operations
 - Overview
 - `std::filesystem::path`
 - File system operations

For brevity, we use namespace stdfs = std::filesystem.

Path

- C++ uses `stdfs::path` to represent a path.
 - It is essentially a string of some *native encoding*, e.g. UTF-8 on Linux, UTF-16 on Windows*.
 - The underlying character can be checked by `stdfs::path::value_type`, normally `char` on Linux and `wchar_t` on Windows.
 - And `stdfs::path::string_type = std::basic_string<value_type>`.
 - And the string stores path in *native format*.
 - You can just access the underlying string in `const` way:

Accesses the native path name as a character string.

```
const value_type* c_str() const noexcept;
const string_type& native() const noexcept;
operator string_type() const;
```

- (1) 1) Equivalent to `native().c_str()`.
(2) 2) Returns the native-format representation of the pathname by reference.
(3) 3) Returns the native-format representation of the pathname by value.

*: Strictly speaking, usually a filesystem doesn't really respect encoding; it just treats the path as some byte sequence (even if it's not a valid UTF-8 / 16). So "native encoding" essentially means "a string that you can pass into filesystem syscall directly", e.g. 2-byte-per-unit sequence on Windows.

Path

Exceptions

2.4-8) May throw implementation-defined exceptions.

- However, you can construct the path by any encoding and format.

```
template< class Source >
path( const Source& source, format fmt = auto_format );
template< class InputIt >
path( InputIt first, InputIt last, format fmt = auto_format );
```

- For **format**, it's essentially a scoped enumeration in **std::fs::path** with three enumerators:

Name	Explanation
native_format	native pathname format
generic_format	generic pathname format
auto_format	implementation-defined format, auto-detected where possible

- By default it uses **auto_format**, i.e. determine if the input format is native or generic automatically and convert if necessary.

Path

```
template< class Source >
path( const Source& source, format fmt = auto_format );
template< class InputIt >
path( InputIt first, InputIt last, format fmt = auto_format );
```

- The template allows you to use any character type, and ctor will convert to its native encoding.
 - (2.1) — `char`: The encoding is the native ordinary encoding. The method of conversion, if any, is operating system dependent.
 - (2.2) — `wchar_t`: The encoding is the native wide encoding. The method of conversion is unspecified.
 - (2.3) — `char8_t`: The encoding is UTF-8. The method of conversion is unspecified.
 - (2.4) — `char16_t`: The encoding is UTF-16. The method of conversion is unspecified.
 - (2.5) — `char32_t`: The encoding is UTF-32. The method of conversion is unspecified.
- For `wchar_t`, Windows won't do any conversion but Linux needs to do so;
- For `char`, Linux won't do any conversion but Windows needs to do so.
- Actually, Windows recognizes `char` for file API in ANSI (or Active) Code Page (ACP).
 - This will lead to complex behavior when interacting with compiler option...

Windows ACP

- Assuming that we have a file path “D:\\试验.txt” on Windows.
 - And we use a default Chinese PC, i.e. ACP is GBK (id: 936).
- Given `main.cpp` as:

```
std::filesystem::path p{ R"(D:\\试验.txt)" };
std::cout << std::filesystem::exists(p) << "\\n";
```
- Case 1: msvc doesn't add any option, and encoding of `main.cpp` is GBK.
 1. D:\\试验.txt is in GBK, and msvc reads it as GBK correctly.
 2. The execution charset is GBK, so D:\\试验.txt is still GBK in binary exe.
 3. Current ACP is GBK, so `stdfs::path` converts it from GBK to native encoding (UTF-16) and stores it;
 4. The path is correct so file system says it exists.

Windows ACP

```
std::filesystem::path p{ R"(D:\试验.txt)" };
std::cout << std::filesystem::exists(p) << "\n";
```

- Case 2: msvc adds `/utf-8`, and encoding of `main.cpp` is UTF-8.
 1. `D:\试验.txt` is in UTF-8, and msvc reads it as UTF-8 correctly.
 2. The execution charset is UTF-8, so `D:\试验.txt` **is UTF-8** in binary exe.
 3. Current ACP is GBK, so `stdfs::path` converts it **from GBK** to native encoding (UTF-16) and stores it;
 - However, UTF-8 string is not really a GBK string, and the corresponding binary leads to GBK as "`D:\璇暭撇.txt`".
 4. The path is not correct and file system says it doesn't exist.
- Case 3: msvc adds `/source-charset:utf-8`, and encoding of `main.cpp` is UTF-8.
 1. `D:\试验.txt` is in UTF-8, and msvc reads it as UTF-8 correctly.
 2. As default execution charset is ACP, `D:\试验.txt` **is GBK** in binary exe.
 3. So the path is still correct and file system says it exists.

Windows ACP

```
std::filesystem::path p{ R"(D:\试验.txt)" };
std::cout << std::filesystem::exists(p) << "\n";
```

- Case 4: msvc adds `/utf-8`, and encoding of `main.cpp` is GBK.
 1. `D:\试验.txt` in GBK is not valid UTF-8, so msvc warns C4828 (illegal character in UTF-8) and silently passes the original bytes as is.
 - So **accidentally**, it's still GBK in binary exe.
 2. Thus accidentally, the path is correct and file system says it exists.
- Case 5: assuming ACP is UTF-8 (65001), msvc doesn't add any option (equiv. to add `/utf-8`), and encoding of `main.cpp` is GBK.
 1. Same as case 4, i.e. the string is of GBK in binary exe.
 2. However, GBK is not valid UTF-8, so ctor of path throws an exception (`std::system_error` in MS-STL).

Path construction

- So to make a valid path, we need first:
 - Make sure the compiler knows how to read the file (string literals), i.e. the file encoding should be correctly specified in source charset.
- And then two ways:
 1. Make execution charset (for string literals) and string encoding (for other strings stored in e.g. `std::string`) same as ACP.
 2. Use `char8_t[] / char16_t[] / char32_t[]` instead.
 - A. Any ACP is OK, since `char8_t` as a unique type will always be decoded as UTF-8 in ctor.
 - B. Any execution charset is OK, since `char8_t` is always UTF-8 in binary exe.

```
std::filesystem::path p{ u8R"(D:\试验.txt)" };
```

- And if you know your `std::string / ...` is essentially UTF-8 while ACP is not, you can `reinterpret_cast` it to avoid conversion.

```
// Assuming we use UTF-8 as execution charset.
std::string s{ R"(D:\试验.txt)" };
auto ptr = reinterpret_cast<const char8_t *>(s.c_str());
std::u8string_view view{ ptr, ptr + s.size() / sizeof(char8_t) };
std::filesystem::path p{ view };
```

Path construction

- On the other hand, on Linux + gcc, no matter what execution charset you use, `char[]` won't do any conversion.
 - As `char` is its native encoding, so libstdc++ assumes correct bytes.
 - Instead, `-fwide-exec-charset` will specify encoding of `wchar_t` and be converted to UTF-8 automatically.
- To specify encoding and conversion explicitly, you can use locale:

```
template< class Source >
path( const Source& source, const std::locale& loc, format fmt = auto_format );


---


template< class InputIt >
path( InputIt first, InputIt last, const std::locale& loc, format fmt = auto_format );
```

- As the conversion is explicitly specified in locale, the template only accepts a char sequence.

Path const

```
path& operator=( const path& p );
path& operator=( path&& p ) noexcept;
path& operator=( string_type&& source );
template< class Source >
path& operator=( const Source& source );
```

```
path& assign( string_type&& source );
template< class Source >
path& assign( const Source& source );
template< class InputIt >
path& assign( InputIt first, InputIt last );
```

- When native encoding is `wchar_t`:
 - Just use `codecvt<wchar_t, char, std::mbstate_t>` in locale to convert `char[]` to native encoding `wchar_t[]`.
 - E.g. For windows, GBK -> UTF-16.
- When native encoding is `char`:
 - First use `codecvt<wchar_t, char, std::mbstate_t>` in locale to convert `char[]` to `wchar_t[]`;
 - E.g. For Linux, GBK -> UTF-32
 - Then convert `wchar_t` back to native encoding `char` (e.g. UTF-8).
 - E.g. UTF-32 -> UTF-8, which is equiv. to construct from `wchar_t[]` directly.
- And finally some omitted overloads & `operator=`, just list here.

```
path() noexcept; Default is just an empty string.
path( const path& p );
path( path&& p ) noexcept;
path( string_type&& source, format fmt = auto_format );
```

Path

- Besides construction, you can also get string of different formats and encodings.
 - Native format + Native encoding: just `.native()`, as we mentioned.
 - Native format + Converted encoding:

```
std::string string() const;
std::wstring wstring() const;
std::u16string u16string() const;
std::u32string u32string() const;
std::u8string u8string() const;
```
 - Generic format + Converted encoding:

```
std::string generic_string() const;
std::wstring generic_wstring() const; (2) (since C++17)
std::u16string generic_u16string() const;
std::u32string generic_u32string() const;
std::u8string generic_u8string() const; (3) (since C++20)
```
- Particularly, these functions are required to use `/` as directory separator.

Path

- And there are also template versions, together will allocator:

```
template< class CharT, class Traits = std::char_traits<CharT>,
          class Alloc = std::allocator<CharT> >
std::basic_string<CharT,Traits,Alloc>
    string( const Alloc& a = Alloc() ) const;
```

```
template< class CharT, class Traits = std::char_traits<CharT>,
          class Alloc = std::allocator<CharT> >
std::basic_string<CharT,Traits,Alloc>
    generic_string( const Alloc& a = Alloc() ) const;
```

- So to get Generic format + Native encoding, just use `.generic_string<stdfs::path::value_type>()`.
- Finally, you can also check the preferred separator by `static constexpr stdfs::path::preferred_separator`.
 - \ on Windows, / on Linux.

Path decomposition

Note: filename with pure extension (e.g. "D:\\\\.gitignore") is not considered as extension (i.e. stem = filename = ".gitignore", extension = "")

- There are also some observer functions to query path components:

- which just return new `stdfs::path`.

`root_name`

`root_directory`

`root_path`

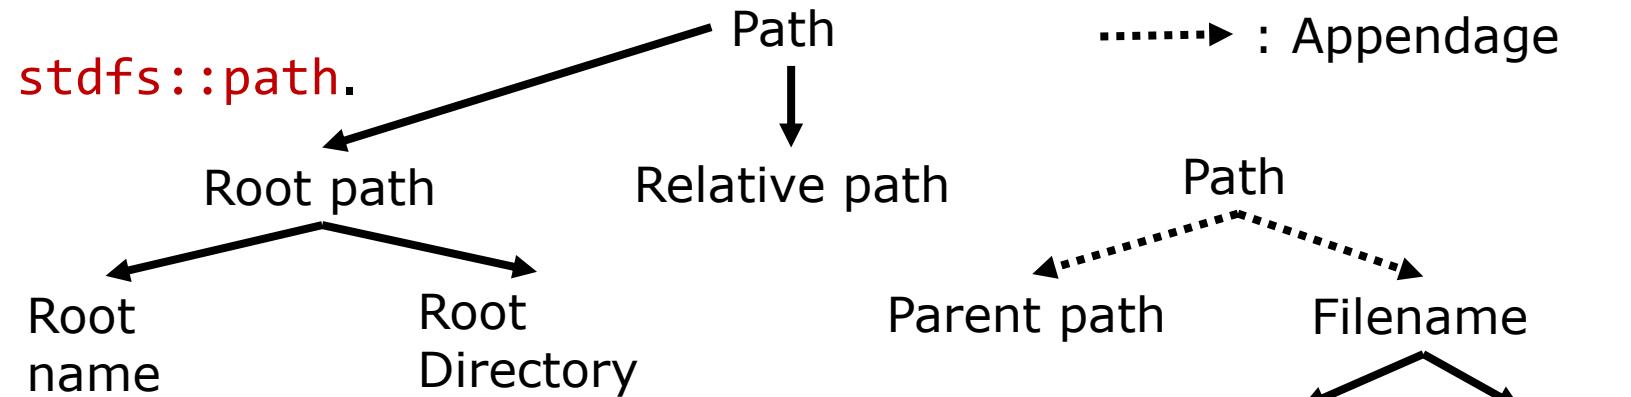
`relative_path`

`parent_path`

`filename`

`stem`

`extension`



```
D:\\sub\\path\\file.txt
Original path = "D:\\\\sub\\\\path\\\\file.txt"
Root name = "D:\\"
Root directory = "\\"
Root path = "D:\\\\"
Relative path = "sub\\\\path\\\\file.txt"
Parent path = "D:\\\\sub\\\\path"
Filename = "file.txt"
Stem = "file"
Extension = ".txt"
```

```
./sub/path/file
Original path = "./sub/path/file"
Root name = ""
Root directory = ""
Root path = ""
Relative path = "./sub/path/file"
Parent path = "./sub/path"
Filename = "file"
Stem = "file"
Extension = ""
```

Path decomposition

- Particularly, such decomposition is **lexical**, which doesn't even really interact with file system.
 - So "parent path" doesn't really return path of parent directory, but just remove the last component.
 - And when a path ends with directory separator, the last component is just empty, so parent just removes the separator.

- For example:

```
D:\sub\path\file\  
Original path = "D:\\sub\\\\path\\\\file\\\\"  
Root name = "D:"  
Root directory = "\\"  
Root path = "D:\\\"  
Relative path = "sub\\\\path\\\\file\\\\"  
Parent path = "D:\\sub\\\\path\\\\file"  
Filename = ""  
Stem = ""  
Extension = ""
```

```
./sub/path/..  
Original path = "./sub/path/.."  
Root name = ""  
Root directory = ""  
Root path = ""  
Relative path = "./sub/path/.."  
Parent path = "./sub/path"  
Filename = ".."  
Stem = ".."  
Extension = ""
```

Note: parent of root directory is still root directory (i.e. "/" -> "/");
but parent of file name is empty (i.e. "data.txt" -> "")).

Path normalization

- To find real parent, you need to do path normalization first.
 - And we say that there are two ways:
 - Lexical: by `.lexically_normal()`; the normalization process is:

Normalization of a generic format pathname means:

1. If the path is empty, stop.

2. Replace each slash character in the *root-name* with a *preferred-separator*.

3. Replace each *directory-separator* with a *preferred-separator*.

Assuming we have a path as
“D:/..\sub\.\Vpath\..\\file.txt”.

D:/..\sub\.\Vpath\..\\file.txt

D:\..\sub\.\Vpath\..\\file.txt

[Note 4: The generic pathname grammar defines *directory-separator* as one or more slashes and *preferred-separators*. — end note]

D:\..\sub\.\path\..\\file.txt

4. Remove each dot filename and any immediately following *directory-separator*.
5. As long as any appear, remove a non-dot-dot filename immediately followed by a *directory-separator* and a dot-dot filename, along with any immediately following *directory-separator*.

D:\..\sub\file.txt

6. If there is a *root-directory*, remove all dot-dot filenames and any *directory-separators* immediately following them.

[Note 5: These dot-dot filenames attempt to refer to nonexistent parent directories. — end note]

7. If the last filename is dot-dot, remove any trailing *directory-separator*. 7. is applied to e.g. ...\\..\\.

8. If the path is empty, add a dot. 1. & 8. An empty path is still empty after normalization, but a non-empty but essentially empty is normalized to ..

Path normalization

- Filesystem-dependent: `stdfs::canonical` / `weakly_canonical`; paths are normalized to a unique absolute path.
 - Path is first converted to an absolute path by `stdfs::absolute(p)`;
 - Then perform lexical normalization.
- `canonical` will check whether the path really exists, while `weakly_canonical` just normalizes it.
- We'll go into details about two forms of global APIs in `stdfs` later.

```
path canonical( const std::filesystem::path& p );
```

(1)

```
path canonical( const std::filesystem::path& p,
                 std::error_code& ec );
```

(2)

```
path weakly_canonical( const std::filesystem::path& p );
```

(3)

```
path weakly_canonical( const std::filesystem::path& p,
                         std::error_code& ec );
```

(4)

Example

```
Original path = "./sub/path/.."  
----- Lexical normal -----  
Path = "sub/"  
Root name = ""  
Root directory = ""  
Root path = ""  
Relative path = "sub/"  
Parent path = "sub"  
Filename = ""  
Stem = ""  
Extension = ""  
----- Weakly canonical -----  
Path = "/app/sub/"  
Root name = ""  
Root directory = "/"  
Root path = "/"  
Relative path = "app/sub/"  
Parent path = "/app/sub"  
Filename = ""  
Stem = ""  
Extension = ""
```

The last / is not stripped, even after normalization. .. only removes **trailing** /.

```
Original path = "./sub/path/"  
----- Lexical normal -----  
Path = "sub/path/"  
Root name = ""  
Root directory = ""  
Root path = ""  
Relative path = "sub/path/"  
Parent path = "sub/path"  
Filename = ""  
Stem = ""  
Extension = ""  
----- Weakly canonical -----  
Path = "/app/sub/path/"  
Root name = ""  
Root directory = "/"  
Root path = "/"  
Relative path = "app/sub/path/"  
Parent path = "/app/sub/path"  
Filename = ""  
Stem = ""  
Extension = ""
```

Path normalization

- Notice that “physical” parent of lexical normalization may still be wrong when normalized result still contains ...
 - Only **(weakly_)canonical** ensures a correct physical parent.
- Finally, you can get or set CWD by **current_path()**:

```
std::cout << std::filesystem::current_path() << "\n";
std::filesystem::current_path("C:\\Temp");
std::cout << std::filesystem::current_path() << "\n";
```

"D:\\Work\\C++\\Tests\\Project1"
"C:\\Temp"

```
Original path = ".../..."
----- Lexical normal -----
Path = ".../..."
Root name = ""
Root directory = ""
Root path = ""
Relative path = ".../..."
Parent path = "..." (highlighted)
Filename = "..."
Stem = "..."
Extension = ""
```

- So **absolute()** is essentially **current_path() / path** when **path** is relative.

DOS directory

```
C:\Users\Public>cd D:\Work  
C:\Users\Public>D:  
D:\Work>C:  
C:\Users\Public>
```

cd: change current directory
“D:” : switch to current directory of Drive D.
“C:” : switch to current directory of Drive C.

- It's also worth noting that DOS maintains separate “current directory” for every drive.
 - So **C:** and **D:** are actually relative paths, while **C:** and **D:** are absolute paths.
- In Windows, current directory is unified as a single path, as known by CWD.
 - However, CMD [pretends](#) they are still there by storing them with “strange environment variables”.
 - And Windows inherits the DOS behavior, regarding **C:** and **D:** as relative paths. But there exists only a single real CWD.

```
std::cout << std::filesystem::current_path() << "\n";  
std::cout << std::filesystem::absolute("C:") << "\n";  
std::cout << std::filesystem::absolute("D:") << "\n";  
std::filesystem::current_path("C:\Temp");  
std::cout << std::filesystem::absolute("C:") << "\n";  
std::cout << std::filesystem::absolute("D:") << "\n";
```

```
"D:\\Work\\C++\\Tests\\Project1"  
"C:\\\"  
"D:\\Work\\C++\\Tests\\Project1"  
"C:\\Temp"  
"D:\\\"
```

Other non-CWD drive will return root.

Path relativization¹

- In contrast to normalization, path can also be “denormalized” by converting an absolute path to relative path.
 - More generally, given a path **b**, how can it be transformed to path **a** with shortest components.
 - For example, `path{ "/a/d" }.relative("/a/b/c")` would be `"../../d"`.
- Similarly, you can use two ways:
 - Lexical: by `a.lexically_relative(b)`; the process is:
 - ① Check whether it's possible to transform **b** to **a**.
 - If it's impossible (i.e. conditions below), return empty path directly.

- `root_name() != base.root_name()` is **true**, or E.g. two paths in different drives in Windows.
- `is_absolute() != base.is_absolute()` is **true**, or Absolute path + relative path, not lexically transformable.
- `!has_root_directory() && base.has_root_directory()` is **true**, or E.g. `bar` and `/foo` on Windows, i.e. you cannot cd from `/foo` to `bar` without CWD.
- any `filename` in `relative_path()` or `base.relative_path()` can be interpreted as a `root-name`,
E.g. for UNC path on Windows, e.g. `\.\C:\Test` uses `C:\Test` as relative path.
That is, UNC doesn't participate in lexical relative process.

Path relativization

```
a = D:\test\test.txt  
b = D:\test\test2\test.txt\..
```

- ② Determine the first mismatched component of two paths (just like `std::mismatch`).

- Let remaining mismatched components of **a** be $[a_1, a_2)$ and **b** be $[b_1, b_2)$:
 - Example above is $[a_1, a_2) = [\text{test}.txt], [b_1, b_2) = [\text{test2}, \text{test}.txt, ..]$.
 - If no mismatched component (i.e. $a_1 = a_2, b_1 = b_2$), return `path{ “.” }`;
 - Otherwise, assuming that in $[b_1, b_2)$, n components are `..` and m components are not `..`, `.` and empty.
 - Example above is $n = 1, m = 2$.
 - If $n > m$ (that is, the lexically normalized form contains only `..`), then return empty path.
 - If $n = m$ (that is, the lexically normalized form is `.`), then:
 - If $a_1 = a_2$, return `path{ “.” }`;
 - Otherwise, return $[a_1, a_2)$.
 - If $n < m$, then return a path with 1. `“..”` repeated for $m - n$ times; 2. $[a_1, a_2)$.

`result = ..\test.txt`

Path relativization

- Actually, this algorithm may falsely report empty path even when such transformation should be possible.

- For example:

```
fs::path p{ "a/b"};
n = 1, m = 0
std::cout << p.lexically_relative("a/b/..") << "\n";
Program returned: 0
"
"
```

- Though theoretically it can be “..”.
- If you want an always-correct lexical transformation, you need to do lexical normalization first.

[Note 3: If normalization (`[fs.path.generic]`) is needed to ensure consistent matching of elements, apply `lexically_normal()` to `*this`, `base`, or both. — end note]

- The second way is filesystem-dependent `stdfs::relative(a, b)`, which will always ensure correct relative path in the file system.
 - It's same as `lexically_relative` two `weakly_canonical` paths.

Path proximation

- Finally there also exists proximation, which means “relativization if possible, otherwise return original path”.

- Effectively:

```
path lexically_proximate(const path& b)
{
    if (auto rel = a.lexically_relative(b); !rel.empty())
        return rel;
    return *this;
}
```

- And `stdfs::proximate(a, b)` is also same as `lexically_proximate` two `weakly_canonical` paths.
- BTW, when `b` is not provided in `relative` and `proximate`, `current_directory` will be used.

```
friend path operator/( const path& lhs, const path& rhs ); path& operator/=( const path& p );
template< class Source >
path& operator/=( const Source& source );
template< class Source >
path& append( const Source& source );
template< class InputIt >
path& append( InputIt first, InputIt last );
```

- For a given path `./sub/path/file.txt`, it's essentially a combination of hierarchical components.
- C++ provides two utilities to combine components.
 1. Append: combine two components with a directory separator, if needed.
 - For example:
`std::filesystem::path p = "/home";`
`// p == /home/tux/.fonts`
`std::cout << p / "tux" / ".fonts" << '\n';`
 - However, there exist lots of corner cases...
 - ① Subpath is absolute path: C++ chooses to overwrite (replace) LHS.
 - For example:
`// On Windows,`
`path("foo") / "C:/bar"; // the result is "C:/bar" (replaces)`
 - But this can be astonishing:
`std::filesystem::path p = "/home";`
`// p == /.fonts`
`std::cout << p / "tux" / "/.fonts" << '\n';`

Reason: `"/.fonts"` is absolute path.

Path composition

- DOS-like behavior on Windows also causes surprising result even when subpath is relative.

- ② No separator is inserted for a single drive; it's still a relative path.

```
std::cout << fs::path{ "C:" } / "Users" / "Admin" << '\n';
std::cout << fs::path{ "C:\\\" } / "Users" / "Admin" << '\n';
```

"C:Users\\Admin"
"C:\\Users\\Admin"

- ③ Appending a relative path that has different drive will replace the whole path;

```
path("foo") / "C:/bar"; // the result is "C:/bar" (replaces)
path("foo") / "C:";    // the result is "C:"      (replaces)
```

- ④ Appending a relative path that has same drive will append as if LHS is CWD.

```
std::cout << fs::path{ "C:\\foo" } / "C:bar" << '\n';
std::cout << fs::path{ "C:foo" } / "C:bar" << '\n';
```

"C:\\foo\\bar"
"C:foo\\bar"

- ⑤ Appending a relative path that has root directory but no drive, to another path with drive will reserve LHS drive.

```
std::cout << fs::path{ "C:\\bar" } / "\\\foo" << '\n';
std::cout << fs::path{ "C:bar" } / "\\\foo" << '\n';
```

"C:\\foo"
"C:\\foo"

Path composition

```
template< class Source >  
path& concat( const Source& source );
```

(7)

```
template< class InputIt >  
path& concat( InputIt first, InputIt last );
```

(8)

2. Concatenate: combine two components as if concatenating underlying strings directly; no additional separator is introduced.

```
path& operator+=( const path& p );
```

(1)

```
path& operator+=( const string_type& str );
```

(2)

```
path& operator+=( std::basic_string_view<value_type> str );
```

```
path& operator+=( const value_type* ptr );
```

(3)

```
path& operator+=( value_type x );
```

(4)

```
template< class CharT >
```

(5)

```
path& operator+=( CharT x );
```

```
template< class Source >
```

(6)

```
path& operator+=( const Source& source );
```

These overloads
are designed to
mimic overloads
of `std::string`
`::operator+=`.

- Strangely, there is no `operator+`; but normally this operation is used to concatenate with a string, so you can just `operator+` all strings first.
 - Or you have to use either `((stdfs::path{a} += b) += c)...` or `.native()` to use `operator+` of `std::basic_string`.

Path composition

- Note 1: Due to associativity, `p / "a" / "b"` is legal while `p /= "a" /= "b"` is illegal.
 - `p / "a" / "b" ⇔ ((p / "a") / "b")`, while
 - `p /= "a" /= "b" ⇔ (p /= ("a" /= "b"))`.
 - And it's illegal for two string literals to `/=`.
 - You have to add a bunch of brackets; that's why we use `((std::path{a} += b) += c)....`
- Note 2: there also exist some boolean observers to check existence.
 - “empty” means the underlying string contains nothing.
 - And `has_xxx` means whether `xxx` is empty or not.

`empty`

`has_root_path`
`has_root_name`
`has_root_directory`
`has_relative_path`
`has_parent_path`
`has_filename`
`has_stem`
`has_extension`

`is_absolute`
`is_relative`

Path iteration

- As a combination of many different components, path can also be iterated (grouped by separator) in generic format.
- It provides `.begin()` and `.end()` that return a path const iterator.
 - It just iterates through root name, root directory and filenames.
 - For example:

```
#include <filesystem>
#include <iostream>
namespace fs = std::filesystem;

int main()
{
    const fs::path p =
#ifdef _WIN32
        "C:\\users\\abcdef\\AppData\\Local\\Temp\\";
#else
        "/home/user/.config/Cppcheck/Cppcheck-GUI.conf";
#endif
    std::cout << "Examining the path " << p << " through iterators gives\n";
    for (auto it = p.begin(); it != p.end(); ++it)
        std::cout << *it << " | ";
    std::cout << '\n';
}
```

```
--- Windows ---
Examining the path "C:\users\abcdef\AppData\Local\Temp\" through iterators gives
"C:" | "/" | "users" | "abcdef" | "AppData" | "Local" | "Temp" | ""

--- UNIX ---
Examining the path "/home/user/.config/Cppcheck/Cppcheck-GUI.conf" through iterators gives
"/" | "home" | "user" | ".config" | "Cppcheck" | "Cppcheck-GUI.conf" |
```

Deferenced result (i.e. `iterator::value_type`) is another `path`.

Path iteration

- Though it seems to be bidirectional iterator, it's actually only regulated to be input iterator.
 - Reason: before C++20, forward iterator has such a regulation:
- If `i` and `j` are both dereferenceable, then `i == j` if and only if `*i` and `*j` are bound to the same object.
 - That is, every component should have a fixed source to make every iterator dereferenced to that source.
 - This requires `path` to store a container of components, making it expensive to construct any `path`...
 - And this is how libstdc++ implements it, making it bidirectional.
- However, path iterator is quite like a string `std::views::split` by separator!
 - Another way (as libc++ and MS-STL do) is to cache the range in the iterator, so only when iterator is used will parsing begins.

Path iteration

- So instead, the standard regulates that:
 - 2 A path::iterator is a constant iterator meeting all the requirements of a bidirectional iterator except that, for dereferenceable iterators a and b of type path::iterator with a == b, there is no requirement that *a and *b are bound to the same object. Its value_type is path.
 - which makes it only satisfies input iterator, and thus it's impossible to apply some functions in `<algorithm>`.
- BUT, such requirement is not part of **bidirectional_iterator** in C++20!
 - So theoretically it should be able to utilize constrained algorithms, i.e. `std::ranges::xxx`.
 - I'm currently planning to submit a DR or proposal to solve this problem (if you're a committee, feel free to contact me ☺).

Path modification

Return reference
to `*this`.

- There also exist some simple non-const methods:

1. `.make_preferred()`: for path whose native format is also generic format, convert current separators to preferred separators.

- For example:

```
fs::path p{ "C:/Test\\Test2" };
std::cout << p << "\n" << p.make_preferred() << "\n";
```

"C:/Test\\Test2"
"C:\\Test\\\\Test2"

2. `.remove_filename()`: Remove the last component (if it exists) so `.has_filename()` returns `false`.

- So after removal, the path is either empty or ends with a separator.

```
std::cout << std::boolalpha
<< (p = "foo/bar").remove_filename()
<< (p = "foo/").remove_filename() <<
<< (p = "/foo").remove_filename() <<
<< (p = "/").remove_filename() << '\'
<< (p = "").remove_filename() << '\t'
```

"foo/"
"foo/"
"/"
"/"
""

“foo” will also be converted to “”.

Modifiers

`clear`

`make_preferred`

`remove_filename`

`replace_filename`

`replace_extension`

`swap`

Path modification

3. `.replace_filename(const path& rep)`: equivalent to 1. `this->remove_filename()`; 2. `(*this) /= rep`.
4. `.replace_extension(const path& rep = {})`: equivalent to code below:
 - For example:

Path:	Ext:	Result:
<code>"/foo/bar.jpg"</code>	<code>".png"</code>	<code>"/foo/bar.png"</code>
<code>"/foo/bar.jpg"</code>	<code>"png"</code>	<code>"/foo/bar.png"</code>
<code>"/foo/bar.jpg"</code>	<code>". "</code>	<code>"/foo/bar."</code>
<code>"/foo/bar.jpg"</code>	<code>" "</code>	<code>"/foo/bar"</code>
<code>"/foo/bar."</code>	<code>"png"</code>	<code>"/foo/bar.png"</code>
<code>"/foo/bar"</code>	<code>".png"</code>	<code>"/foo/bar.png"</code>
<code>"/foo/bar"</code>	<code>"png"</code>	<code>"/foo/bar.png"</code>
<code>"/foo/bar"</code>	<code>". "</code>	<code>"/foo/bar."</code>
<code>"/foo/bar"</code>	<code>" "</code>	<code>"/foo/bar"</code>
<code>"/foo/.."</code>	<code>".png"</code>	<code>"/foo/..png"</code>
<code>"/foo/.."</code>	<code>"png"</code>	<code>"/foo/..png"</code>
<code>"/foo/.."</code>	<code>". "</code>	<code>"/foo/.."</code>
<code>"/foo/.."</code>	<code>" "</code>	<code>"/foo/.."</code>
<code>"/foo/"</code>	<code>".png"</code>	<code>"/foo/.png"</code>
<code>"/foo/"</code>	<code>"png"</code>	<code>"/foo/.png"</code>

```
// Assume we have a function TryRemoveExtension,  
// which will remove extension if it exists.  
path& ReplaceExtension(const path& rep = {})  
{  
    TryRemoveExtension();  
    if (rep.empty())  
        return *this;  
    // Add '.' if necessary  
    if (rep.native()[0] != DOT)  
        this->concat(DOT);  
    return *this += rep;  
}
```

Path

- And finally some simple utilities, just list here.
 - Comparable (by `<=>/==` or `.compare`); compare **components**. `== "D://Test"`
 - Hashable (by `std::hash` or `friend hash_value`); hash components.
 - Input / Output by `>>` / `<<`;
 - For `iostream<CharT, Traits>`, equiv. to input / output the `.string <CharT, Traits>()` with `std::quoted` so space will not interrupt input. `"C:\\foo\\bar"`
 - Recap: `quoted` will escape the quote and the escape, so `\` is escaped to `\\`.
 - Formattable **since C++26**.

Format specification

The syntax of format specifications *path-format-spec* is:

fill-and-align(optional) *width*(optional) ?(optional) **g**(optional)

fill-and-align and *width* have the same meaning as in standard format specification.

The **?** option is used to format the pathname as an escaped string.

The **g** option is used to specify that the pathname is in generic-format representation.

Path formatting

- Its regulation is slightly different from `.string()`.

`std::formatter<std::filesystem::path>::format`

```
template< class FormatContext >
auto format( const std::filesystem::path& p, FormatContext& ctx ) const
    -> FormatContext::iterator;
```

Let `s` be `p.generic_string<std::filesystem::path::value_type>()` if the `g` option is used, otherwise `p.native()`. Writes `s` into `ctx.out()` as specified by *path-format-spec*.

For character transcoding of the pathname:

- The pathname is transcoded from the native encoding for wide character strings to UTF-8 with maximal subparts of ill-formed subsequences substituted with U+FFFD REPLACEMENT CHARACTER if
 - `std::is_same_v<CharT, char>` is `true`,
 - `std::is_same_v<typename path::value_type, wchar_t>` is `true`, and
 - ordinary literal encoding is UTF-8.
- Otherwise, no transcoding is performed if `std::is_same_v<typename path::value_type, CharT>` is `true`.
- Otherwise, transcoding is implementation-defined.

Returns an iterator past the end of the output range.

i.e. explicitly regulate that in UTF-8 execution charset on Windows, illegal character will be converted to U+FFFD. 

As a C++26 feature, it's not yet implemented so "implementation-defined" is unknown (but likely to be printable with `std::print`).

i.e. encoding of char string literal, as specified in compiler execution charset option. 

Final Notes

- Note 1: there also exists `.u8path()` to construct from a UTF-8 string in C++17, which is then deprecated in C++20.
 - Reason: C++20 introduces `char8_t`, which distinguishes UTF-8 sequence from `char` by template so there is no need to introduce a new method.
 - For the same reason, `.(generic_)u8string` returns `std::string` in C++17 and `std::u8string` in C++20.
- Note 2: to use `path` as key in map, you usually need to normalize it first by `canonical`.
 - Reason: comparison and hashing of `path` are performed **lexically** for the underlying components.
 - Without normalization, two equivalent paths may be seen as two keys.
 - On Windows, you may even need to `to_lower` all case-insensitive paths.
 - A more expensive but always-correct way is by `stdfs::equivalent` to compare, which needs file system call to check equality of two paths.

Covered later.